

Exceptions are the Norm

Dealing with Bad Actors in ETL

Sameer Agarwal

Spark Summit | Boston | Feb 9th 2017

About Me

- Software Engineer at Databricks (Spark Core/SQL)
- PhD in Databases (AMPLab, UC Berkeley)
- Research on BlinkDB (Approximate Queries in Spark)



Overview

1. What's an ETL Pipeline?

- How is it different from a regular query execution pipeline?

2. Using SparkSQL for ETL

- Dealing with Dirty Data (Bad Records or Files)
- Performance (Project Tungsten)

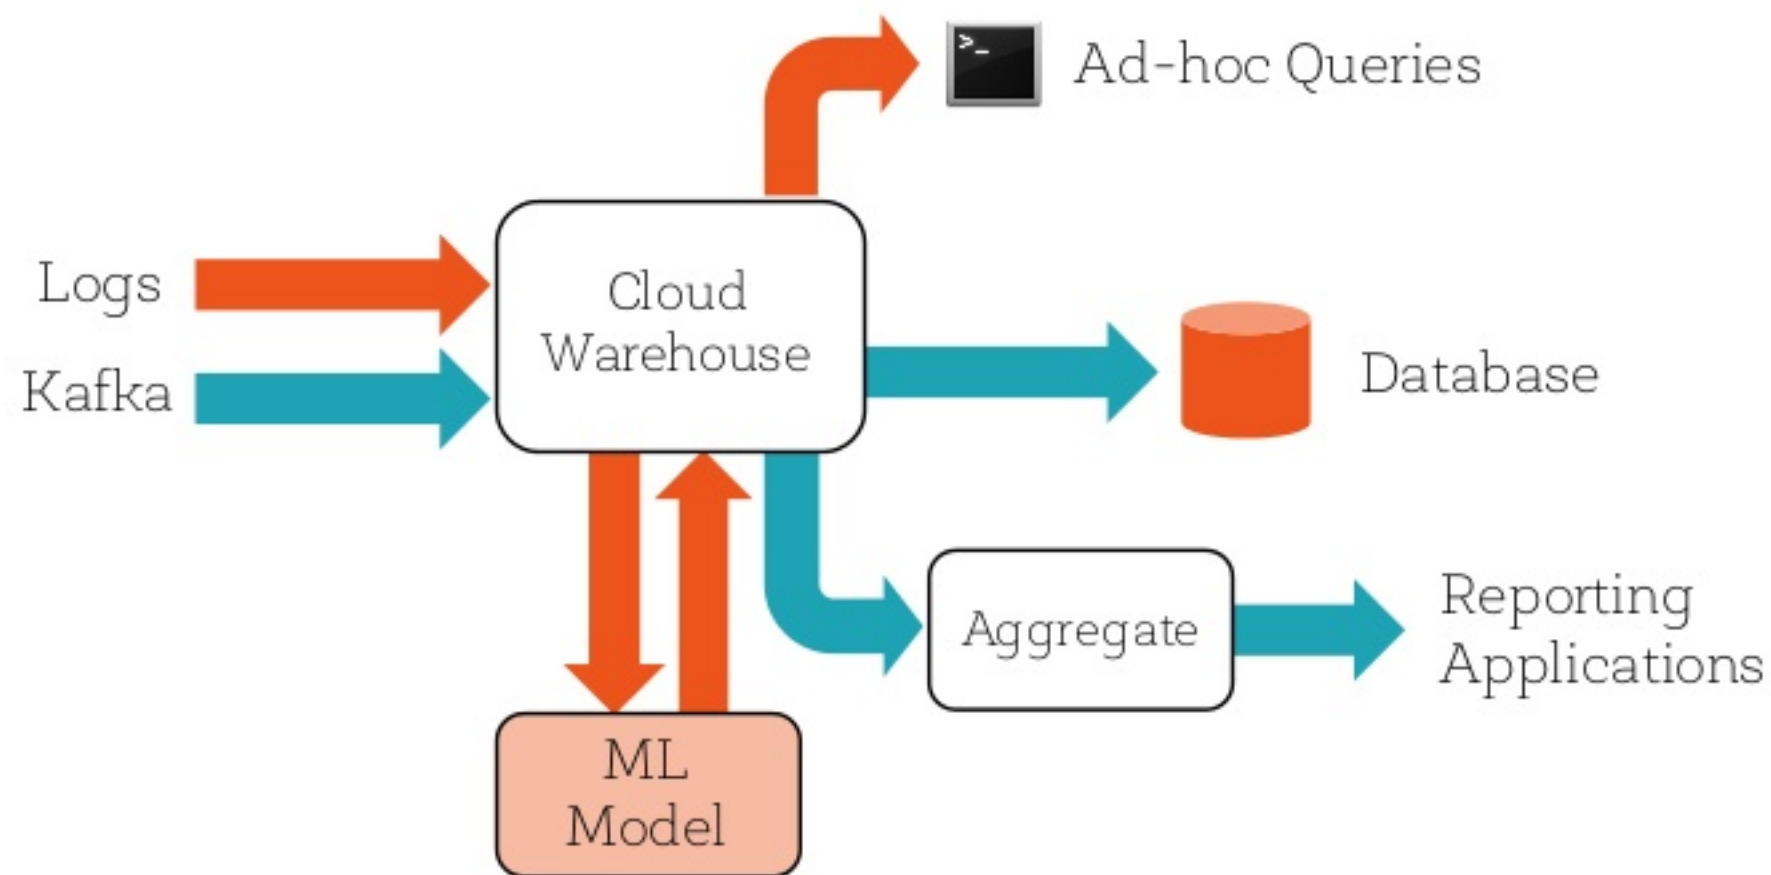
3. New Features in Spark 2.2 and 2.3

- Focus on building ETL-friendly pipelines

What is a Data Pipeline?

1. Sequence of transformations on data
2. Source data is typically semi-structured/unstructured (JSON, CSV etc.)
3. Output data is structured and ready for use by analysts and data scientists
4. Source and destination are often on different storage systems.

Example of a Data Pipeline



ETL is the First Step in a Data Pipeline

1. ETL stands for EXTRACT, TRANSFORM and LOAD
2. Goal is to “clean” or “curate” the data
 - Retrieve data from source (EXTRACT)
 - Transform data into a consumable format (TRANSFORM)
 - Transmit data to downstream consumers (LOAD)

An ETL Query in Spark

```
spark.read.csv("/source/path")
```

EXTRACT

An ETL Query in Spark

```
spark.read.csv("/source/path")
```

EXTRACT

```
.filter(...)
```

```
.agg(...)
```

TRANSFORM

An ETL Query in Spark

```
spark.read.csv("/source/path")  
  .filter(...)  
  .agg(...)
```

EXTRACT

TRANSFORM

```
  .write.mode("append")  
  .parquet("/output/path")
```

LOAD

What's so hard about ETL
Queries?

Why is ETL Hard?

1. Data can be Messy

- Incomplete information
- Missing data stored as empty strings, "none", "missing", "xxx" etc.

2. Data can be Inconsistent

- Data conversion and type validation in many cases is error-prone
 - For e.g., expecting a number but found "123 000"
 - different formats "31/12/2017" "12/31/2017"
- Incorrect information
 - For e.g., expecting 5 fields in CSV, but can't find 5 fields.

Why is ETL Hard?

3. Data can be Constantly Arriving

- At least once or exactly once semantics
- Fault tolerance
- Scalability

4. Data can be Complex

- For e.g., Nested JSON data to extract and flatten
- Dealing with inconsistency is even worse

This is why ETL is important

Consumers of this data don't want to deal with this messiness and complexity

On the flip side


1. A few bad records can fail a job
 - These are not the same as transient errors
 - No recourse for recovery
2. Support for ETL features
 - File formats and conversions have gaps
 - For e.g., multi-line support, date conversions
3. Performance

Spark's **flexible APIs**, support for a wide variety of **datasources** and **state of art tungsten execution engine** makes it a great framework for building end-to-end ETL Pipelines

Using SparkSQL for ETL


Dealing with Bad Data: Skip Corrupt Files

```
spark.read.csv("/source/path")  
  .filter(...)  
  .agg(...)  
  .write.mode("append")  
  .parquet("/output/path")
```




Dealing with Bad Data: Skip Corrupt Files

```
spark.read.csv("/source/path")  
  .filter(...)  
  .agg(...)  
  .write.mode("append")  
  .parquet("/output/path")
```




Dealing with Bad Data: Skip Corrupt Files

[SPARK-17850] If true, the Spark jobs will continue to run even when it encounters corrupt or non-existent files. The contents that have been read will still be returned.




Dealing with Bad Data: Skip Corrupt Records



Dealing with Bad Data: Skip Corrupt Records

[SPARK-12833][SPARK-13764] TextFile formats (JSON and CSV) support 3 different ParseModes while reading data:

1. PERMISSIVE
2. DROPMALFORMED
3. FAILFAST



JSON: Dealing with Corrupt Records

Can be configured via


```
{ "a": 1, "b": 2, "c": 3 }
```

```
{ "a": {, b: 3 }
```

```
{ "a": 5, "b": 6, "c": 7 }
```

`spark.sql.columnNameOfCorruptRecord`

```
spark.read  
  .option("mode", "PERMISSIVE")  
  .json(corruptRecords)  
  .show()
```



<code>_corrupt_record</code>	a	b	c
null	1	2	3
<code>{"a":{, b:3}</code>	null	null	null
null	5	6	7

JSON: Dealing with Corrupt Records

```
{ "a": 1, "b": 2, "c": 3 }
```

```
{ "a": {, b: 3 }
```

```
{ "a": 5, "b": 6, "c": 7 }
```

```
spark.read  
  .option("mode", "DROPMALFORMED")  
  .json(corruptRecords)  
  .show()
```

a	b	c	
1	2	3	
5	6	7	

JSON: Dealing with Corrupt Records

```
{ "a": 1, "b": 2, "c": 3 }  
{ "a": {, b: 3 }  
{ "a": 5, "b": 6, "c": 7 }
```

```
spark.read  
  .option("mode", "FAILFAST")  
  .json(corruptRecords)  
  .show()
```

```
org.apache.spark.sql.catalyst.json  
  .SparkSQLJsonProcessingException:  
Malformed line in FAILFAST mode:  
{"a":{, b:3}
```

CSV: Dealing with Corrupt Records

```
year,make,model,comment,blank
"2012","Tesla","S","No comment",
1997,Ford,E350,"Go get one now they",
2015,Chevy,Volt
```

```
spark.read
  .format("csv")
  .option("mode", "PERMISSIVE")
  .load(corruptRecords)
  .show()
```

year	make	model	comment	blank
2012	Tesla	S	No comment	null
1997	Ford	E350	Go get one now th...	null
2015	Chevy	Volt	null	null

CSV: Dealing with Corrupt Records

```
year,make,model,comment,blank
"2012","Tesla","S","No comment",
1997,Ford,E350,"Go get one now they",
2015,Chevy,Volt
```

```
spark.read
  .format("csv")
  .option("mode", "DROPMALFORMED")
  .load(corruptRecords)
  .show()
```

year	make	model	comment	blank
2012	Tesla	S	No comment	null
1997	Ford	E350	Go get one now th...	null

CSV: Dealing with Corrupt Records


```
year,make,model,comment,blank  
"2012","Tesla","S","No comment",  
1997,Ford,E350,"Go get one now they",  
2015,Chevy,Volt
```

```
spark.read  
  .format("csv")  
  .option("mode", "FAILFAST")  
  .load(corruptRecords)  
  .show()
```

```
java.lang.RuntimeException:  
Malformed line in FAILFAST mode:  
2015,Chevy,Volt
```

Spark Performance: Project Tungsten

Substantially improve the memory and CPU efficiency of Spark **backend execution** and push performance closer to the limits of modern hardware.



Spark Performance: Project Tungsten

Phase 1

Foundation

Memory Management

Code Generation

Cache-aware Algorithms

Phase 2

Order-of-magnitude Faster

Whole-stage Codegen

Vectorization

SparkSQL: A Compiler from Queries to RDDs (Developer Track at 5:40pm)

Operator Benchmarks: Cost/Row (ns)

primitive	Spark 1.6	Spark 2.0
filter	15 ns	1.1 ns
sum w/o group	14 ns	0.9 ns
sum w/ group	79 ns	10.7 ns
hash join	115 ns	4.0 ns
sort (8-bit entropy)	620 ns	5.3 ns
sort (64-bit entropy)	620 ns	40 ns
sort-merge join	750 ns	700 ns
Parquet decoding (single int column)	120 ns	13 ns

**5-30x
Speedups**

Operator Benchmarks: Cost/Row (ns)

primitive	Spark 1.6	Spark 2.0
filter	15 ns	1.1 ns
sum w/o group	14 ns	0.9 ns
sum w/ group	79 ns	10.7 ns
hash join	115 ns	4.0 ns
sort (8-bit entropy)	620 ns	5.3 ns
sort (64-bit entropy)	620 ns	40 ns
sort-merge join	750 ns	700 ns
Parquet decoding (single int column)	120 ns	13 ns

Radix Sort
10-100x
Speedups

Operator Benchmarks: Cost/Row (ns)

primitive	Spark 1.6	Spark 2.0
filter	15 ns	1.1 ns
sum w/o group	14 ns	0.9 ns
sum w/ group	79 ns	10.7 ns
hash join	115 ns	4.0 ns
sort (8-bit entropy)	620 ns	5.3 ns
sort (64-bit entropy)	620 ns	40 ns
sort-merge join	750 ns	700 ns
Parquet decoding (single int column)	120 ns	13 ns


Shuffling
still the
bottleneck

Operator Benchmarks: Cost/Row (ns)

primitive	Spark 1.6	Spark 2.0
filter	15 ns	1.1 ns
sum w/o group	14 ns	0.9 ns
sum w/ group	79 ns	10.7 ns
hash join	115 ns	4.0 ns
sort (8-bit entropy)	620 ns	5.3 ns
sort (64-bit entropy)	620 ns	40 ns
sort-merge join	750 ns	700 ns
Parquet decoding (single int column)	120 ns	13 ns

**10x
Speedup**

TPC-DS (Scale Factor 1500, 100 cores)



Apache Spark 2.2 and 2.3

Massive focus on building ETL-friendly pipelines

New Features in Spark 2.2 and 2.3

1. Better Functionality:

- Improved JSON and CSV Support

2. Better Usability:

- Better Error Messages

3. Better Performance:

- SQL Execution
- Python UDF Processing

Functionality: Better JSON Support

1. [[SPARK-18352](#)] Multi-line JSON Support
 - Spark currently reads JSON one line at a time
 - This currently requires custom ETL

```
spark.read  
  .option("wholeFile",true)  
  .json(path)
```

Functionality: Better JSON Support

2. [[SPARK-19480](#)] Higher order functions in SQL
 - Enable users to manipulate nested data in Spark
 - Operations include map, filter, reduce on arrays/maps

```
tbl_x
```

```
|-- key: long (nullable = false)
|-- values: array (nullable = false)
|    |-- element: long (containsNull = false)
```

Functionality: Better JSON Support

2. [[SPARK-19480](#)] Higher order functions in SQL

```
tbl_x
```

```
|-- key: long (nullable = false)
```

```
|-- values: array (nullable = false)
```

```
|    |-- element: long (containsNull = false)
```

```
SELECT key, TRANSFORM(values, v -> v + key)
FROM tbl_x
```

Functionality: Better CSV Support

1. [[SPARK-16099](#)] Improved/Performant CSV Datasource
 - Multiline CSV Support
 - Additional options for CSV Parsing
 - Whole text reader for dataframes

Functionality: Better ETL Support

1. More Fine-grained (record-level) tolerance to errors
 - Provide users with controls on how to handle these errors
 - Ignore and report errors post-hoc
 - Ignore bad rows up to a certain number or percentage

Usability: Better Error Messages

1. Spark must explain why data is bad
2. This is especially true for data conversion
 - `scala.MatchError: start (of class java.lang.String)`
3. Which row in your source data could not be converted ?
4. Which column could not be converted ?

Performance: SQL Execution

1. SPARK-16026: Cost Based Optimizer
 - Leverage table/column level statistics to optimize joins and aggregates
 - Statistics Collection Framework (Spark 2.1)
 - Cost Based Optimizer (Spark 2.2)
2. Boosting Spark's Performance on Many-Core Machines
 - In-memory/ single node shuffle
3. Improving quality of generated code and better integration with the in-memory column format in Spark

Performance: Python UDFs

1. Python is the most popular language for ETL
2. Python UDFs are often used to express elaborate data conversions/transformations
3. Any improvements to python UDF processing will ultimately improve ETL.
4. **Next talk:** Improving Python and Spark Performance and Interoperability (Wes McKinney)

Recap

1. What's an ETL Pipeline?

- How is it different from a regular query execution pipeline?

2. Using SparkSQL for ETL

- Dealing with Dirty Data (Bad Records or Files)
- Performance (Project Tungsten)

3. New Features in Spark 2.2 and 2.3

- Focus on building ETL-friendly pipelines

Questions?